

Aim : To perform Simple Linear Regression and find out the coefficients of it.

```
In [1]: #Name : Rajshri Kirandas Satpute
#Roll No. : 55
#Year : 3rd year
#Section : B
#Date :09/10/2023

In [2]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

In [3]: import os

In [4]: os.getcwd()

Out[4]: 'C:\\Users\\fatin\\Downloads'

In [5]: os.chdir("C:\\Users\\fatin\\OneDrive\\Desktop")

In [5]: df=pd.read_csv("Salary_Data.csv")

In [6]: df.head()

Out[6]:
   YearsExperience  Salary
0              1.1   39343.0
1              1.3   46205.0
2              1.5   37731.0
3              2.0   43525.0
4              2.2   39891.0

In [7]: df.tail()

Out[7]:
   YearsExperience  Salary
25              9.0  105582.0
26              9.5  116969.0
27              9.6  112635.0
28             10.3  122391.0
29             10.5  121872.0

In [8]: df.head(30)

Out[8]:
   YearsExperience  Salary
0              1.1   39343.0
1              1.3   46205.0
2              1.5   37731.0
3              2.0   43525.0
4              2.2   39891.0
5              2.9   56642.0
6              3.0   60150.0
7              3.2   54445.0
8              3.2   64445.0
9              3.7   57189.0
10             3.9   63218.0
11             4.0   55794.0
12             4.0   56957.0
13             4.1   57081.0
14             4.5   61111.0
15             4.9   67938.0
16             5.1   66029.0
17             5.3   83088.0
18             5.9   81363.0
19             6.0   93940.0
20             6.8   91738.0
21             7.1   98273.0
22             7.9  101302.0
23             8.2  113812.0
24             8.7  109431.0
25             9.0  105582.0
26             9.5  116969.0
27             9.6  112635.0
28            10.3  122391.0
29            10.5  121872.0

In [9]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
 #   Column        Non-Null Count  Dtype
---  --
 0   YearsExperience  30 non-null    float64
 1   Salary          30 non-null    float64
dtypes: float64(2)
memory usage: 608.0 bytes

In [10]: df.describe()

Out[10]:
   YearsExperience  Salary
count            30.000000    30.000000
mean             5.313333    76003.000000
std              2.837888    27414.429785
min              1.100000    37731.000000
25%              3.200000    56720.750000
50%              4.700000    65237.000000
75%              7.700000   100544.750000
max             10.500000   122391.000000

In [11]: df.shape

Out[11]: (30, 2)

In [12]: df.size

Out[12]: 60

In [13]: df.ndim

Out[13]: 2

In [14]: df.isnull().sum()

Out[14]:
YearsExperience    0
Salary            0
dtype: int64

In [15]: #Assiging values in X & Y
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values

#X = df['YearsExperience']
#y = df['Salary']

In [16]: print(X)

[[ 1.1]
 [ 1.3]
 [ 1.5]
 [ 2. ]
 [ 2.2]
 [ 2.9]
 [ 3. ]
 [ 3.2]
 [ 3.2]
 [ 3.7]
 [ 3.9]
 [ 4. ]
 [ 4. ]
 [ 4.1]
 [ 4.5]
 [ 4.9]
 [ 5.1]
 [ 5.3]
 [ 5.9]
 [ 6. ]
 [ 6.8]
 [ 7.1]
 [ 7.9]
 [ 8.2]
 [ 8.7]
 [ 9. ]
 [ 9.5]
 [ 9.6]
[10.3]
[10.5]]

In [17]: print(y)

[ 39343.  46205.  37731.  43525.  39891.  56642.  60150.  54445.  64445.
 57189.  63218.  55794.  56957.  57081.  61111.  67938.  66029.  83088.
 81363.  93940.  91738.  98273. 101302. 113812. 109431. 105582. 116969.
112635. 122391. 121872.]

In [18]: #Splitting testdata into X_train,X_test,y_train,y_test
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=.3,random_state=42)

In [19]: print(X_train)

[[ 1.1]
 [ 2.2]
 [ 5.1]
 [ 2.9]
 [ 4.1]
 [ 4. ]
 [ 7.9]
 [ 1.3]
 [ 1.5]
 [ 9. ]
 [ 2. ]
 [ 7.1]
 [ 9.5]
 [ 5.9]
[10.5]
 [ 6.8]
 [ 3.2]
 [ 3.9]
 [ 4.5]
 [ 6. ]
 [ 3. ]]

In [20]: print(X_test)

[[ 9.6]
 [ 4.9]
 [ 8.2]
 [ 5.3]
 [ 3.2]
 [ 3.7]
[10.3]
 [ 8.7]
 [ 4. ]]

In [21]: print(y_train)

[ 39343.  39891.  66029.  56642.  57081.  55794. 101302.  46205.  37731.
105582.  43525.  98273. 116969.  81363. 121872.  91738.  54445.  63218.
 61111.  93940.  60150.]

In [22]: print (y_test)

[112635.  67938. 113812.  83088.  64445.  57189. 122391. 109431.  56957.]

In [23]: from sklearn.linear_model import LinearRegression
lr = LinearRegression()
lr.fit(X_train, y_train)

Out[23]: LinearRegression()

In [24]: #Assigning Coefficient (slope) to m
m = lr.coef_

In [25]: print("Coefficient :", m)

Coefficient : [9339.08172382]

In [26]: #Assigning Y-intercept to a
c = lr.intercept_

In [27]: print("Intercept :", c)

Intercept : 25918.438334893202

In [ ]:
```